

ING5SE – 2022

Project : connected components labelling

Etienne Hamelin (ehamelin@omnesintervenant.com),

Alexandre Berne (aberne@omnesintervenant.com)

1 Introduction

This project is like a bigger lab. It will require to dig deeper into algorithmic aspects ; not only programming language. Take the time to understand, then solve, the algorithmic problem. Then, programming the solution becomes trivial.

2 Presentation

You will work on connected components labelling, an algorithm often used in computer vision and robotics. The purpose is to distinguish and count, objects in a picture. In the following black/white image, it is obvious to a human eye to distinguish every single bolt, and counting them is (more-or-less) simple. To a computer however, distinguishing them requires some non-trivial work.




Figure 1 – binary image : bolts



Figure 2 – color output : every single bolt identified

2.1 Reminder: images & computer vision

An image of width W and height H is an array or dimension $W \times H$, where each element $I_{x,y}$ is the color of the pixel (x,y) . For historical reasons, in computer systems the origin, i.e. coordinates $(0,0)$, is usually located on the top-left corner, with the y axis pointing down ; the bottom-right corner is at coordinates $(W - 1, H - 1)$.




Figure 3 – image pixels representation

2.1.1 Proposed library

We provide you with a small library, that supports reading and writing NetBPM image files. You should use the `display` command to view images, and `convert` to transform to/from other image formats, both installed with the `imagemagick` package.

This library supports several image types :

- Black or white, or binary images (type=IMAGE_BITMAP),
- Grayscale images, with gray levels encoded either with 8 or 16 bits (IMAGE_GRAYSCALE_8 and IMAGE_GRAYSCALE_16)
- Color images, where the red, green, blue, components of each pixel are encoded with 8 bits in that order (IMAGE_RGB_888).

2.1.2 Main library APIs

The example code below gives is a reference on how to use the library.

```
/* Create an object for a 320 (width) x 200 (height) binary image */
image_t *img = image_new(320, 200, IMAGE_BITMAP);

/* read, write, the binary color of a pixel at (x,y) */
bool b = image_bmp_getpixel(img, x, y).bit;
image_bmp_setpixel(img, x, y, (color_t){.bit = 1});

/* Create a color image */
image_t *img = image_new(320, 200, IMAGE_RGB_888);

/* Read, write the color of a RGB_888 pixel */
uint8_t r = image_rgb_getpixel(img, x, y).rgb.r;
image_rgb_setpixel(img, x, y,
    (color_t){.rgb = {.r = 255, .g = 255, .b = 0}});

/* Read the dimensions of an image object */
int w = img->width;
int h = img->height;

/* Check is coordinates (x,y) are within (0,0)...(w-1,h-1) range */
image_coord_check(img,x,y)

/* Read, write a NetBPM image file */
image_t *img = image_new_open("file.ppm");
image_save_binary(img, "file.ppm")
```

3 Connex components labelling

For simplification, we take as input a binary black/white image, where white (foreground) indicates the presence of an object, and black is the background. Connex components labelling consists in giving a number (a tag) to each white pixel, so that all adjacent foreground pixels share the same tag.

3.1 Definitions

3.1.1 Connectedness

Pixels $p_A = (x_A, y_A)$ and $p_B = (x_B, y_B)$ are said to be adjacent iff $\begin{cases} x_A = x_B \text{ and } y_A = y_B \pm 1 \\ \text{or} \\ x_A = x_B \pm 1 \text{ and } y_A = y_B \end{cases}$.

We usually denote with north, south, east, west (N/S/E/W) the 4 adjacent pixels of a given pixel.



Figure 4 – Yellow pixels are adjacent to the red one

3.1.2 Connex component

A connex component \mathcal{A} of image I is a set of pixels so that for each pair of pixels $p_A, p_B \in \mathcal{A}$, there exists a continuous path of adjacent pixels from p_A to p_B that share the same color in image I .

4 The Rosenfeld & Pfalz algorithm

The Rosenfeld & Pfalz algorithm (proposed en 1966) is one of the simplest algorithms for connex components labelling. It runs in 3 steps:

1. First pass, mark pixels with a temporary tag,

At that stage, a single connex component might bear several distinct tags; but we note in a table when several tags are linked to the same connected component: an equivalence table

2. Analyze the equivalence table, in order to assign a definitive class tag to every temporary tag
3. Replace temporary tags with the definitive class tag.

4.1 Algorithm definition

4.1.1 Algorithm: initial marking

Algorithm	Comments
<p>Inputs : image I, of size $W \times H$</p> <p>Outputs :</p> <ul style="list-style-type: none"> - image E of temporary tags (size $W \times H$), - number of temporary tags n_E, - equivalence table T (of size n_E^{max}) <p>Initialization :</p> <ul style="list-style-type: none"> - $n_E \leftarrow 0$ - $T \leftarrow [0, ..., 0]$ - $E \leftarrow [[0, ... 0], ..., [0, ..., 0]]$ <p>Procedure :</p> <ul style="list-style-type: none"> - For y from 0 to $H - 1$: <ul style="list-style-type: none"> o For x from 0 to $W - 1$: <ul style="list-style-type: none"> ▪ if $I_{x,y} = 0$: 	 <p><i>For each pixel:</i></p> <p><i>Background ? tag = 0</i></p>

<ul style="list-style-type: none"> • $E_{x,y} \leftarrow 0$ ▪ else : <ul style="list-style-type: none"> • $e_N = E_{x,y-1}$ • $e_W = E_{x-1,y}$ • if $e_N = 0$ and $e_W = 0$: <ul style="list-style-type: none"> ◦ $n_E \leftarrow n_E + 1$ ◦ $T[n_E] \leftarrow n_E$ ◦ $E_{x,y} \leftarrow n_E$ • else: <ul style="list-style-type: none"> ◦ $E_{x,y} = \minNonZero(e_N, e_W)$ • if $e_N > 0$ and $e_W > 0$ and $e_N \neq e_W$: <ul style="list-style-type: none"> ◦ $Union(T, e_N, e_W)$ 	<p><i>Foreground ?</i> <i>Read the temp. tag of north/east pixels (already processed)</i></p> <p><i>N/E neighbors not yet tagged?</i> <i>Create new tag</i> <i>With no equivalence set</i></p> <p><i>N or E is already tagged? Set the minimal non-zero tag</i></p> <p><i>N and E have different tags? Note that these tags are now equivalent.</i></p>
--	--

4.1.2 Algorithm: Find

Algorithm	Comments
Inputs : equivalence table T , tag e Outputs : r , the root tag of equivalence class that contains tag e Initialization : - $r \leftarrow e$ Procedure : - while $T[r] < r$: ◦ $r \leftarrow T[r]$	<p>Note that this recursion terminates always, since by construction $\forall e$, $T[e] \leq e$.</p>

4.1.3 Algorithm: Union

Algorithm	Comments
Inputs : equivalence table T , two tags e_1 and e_2 Outputs : modified equivalence table T Initialization : - $r_1 \leftarrow Find(T, e_1)$ - $r_2 \leftarrow Find(T, e_2)$ Procedure : - if $r_1 < r_2$: ▪ $T[r_2] = r_1$ - else : ▪ $T[r_1] = r_2$	<p>Find the root of the equivalence classes containing tags e_1 and e_2</p> <p>Mark equivalence $r_2 \sim r_1$</p>

4.1.4 Algorithm: Renum

Algorithm	Comments
Inputs : table T Outputs : table N ; number of connected components n_c Initialization : - $N \leftarrow [0, \dots, 0]$ of size n_E^{max} - $n_c \leftarrow 0$ Procedure : - For e from 1 to n_E : ◦ if $T[e] = e$: ▪ $n_c \leftarrow n_c + 1$ ▪ $N[e] \leftarrow n_c$ ◦ else : ▪ $N[e] \leftarrow N[T[e]]$	<p>If $T[e] = e$: tag e is the root of an equivalence class; create a new equivalence class number (definitive tag).</p>

	Otherwise: tag e is the child of $T[e]$, and should use the same definitive tag
--	--

5 Provided implementation

You will find attached a correct sequential implementation of the Rosenfeld&Pfalz algorithm

5.1 Useful commands

Clean up temporary files: `make clean`

Run in debug mode: the code prints out intermediate results, and generates additional visualization files.

```
make clean && make DEBUG=1
./main img/test1.pbm
```

Visualize intermediate data outputs:

```
code tags.pgm          # read temporary tags
code classes.pgm       # read definitive tags
display color.ppm      # read the color-coded tags
```

Run in performance mode (`DEBUG=0`): the code prints out less detail, and does not generate color images or temporary tags→ *use this mode to measure execution times*.

```
make clean && make DEBUG=0
./main img/test1.pbm
```

If you encounter memory problems:

```
valgrind ./main img/test1.pbm
```

If you need to convert a ppm (color), pgm (grayscale) or pbm (binary) image to/from another file format:

```
convert file.ppm file.png
```

To convert an image of yours to binary black/white:

```
convert file.png -threshold 50% file.pbm
```

5.2 Code structure

```
|— img          a few test files
|   |— boulons.pbm  a few “larger” images to test performances
|   |— cadastre.pbm idem
|   |— ocr.pbm      idem
|   |— test0.pbm    smalled examples, to understand how it works (or doesn’t)
|   |— test1.pbm
|— inc          header files
|   |— image_connected_components.h
|   |— image_file_io.h
|   |— image.h
```

```

|   |— image_lib.h
|   |— pixel.h
|   |— utils.h
|— Makefile
|— src          source code
    |— main.c    launch CCL algorithm
    |— image.c    basic image manipulation
    |— image_connected_components.c the main CCL algorithm
    |— image_file_io.c read/write NetBPM files
    |— pixel.c    pixel color format conversions

```

Please only change the `image_connected_components.c` and `main.c` files.

Here is the main sequence of functions:

```

main() [main.c]
|— test_image_connected_components [main.c]
    |— image_connected_components [image_connected_components.c]
        |— ccl_temp_tag Initial marking algorithm
        |— ccl_reduce_equivalences renum algorithm
        |— ccl_retag Definitive tag marking
        |— ccl_analyze Analyze each connected component
        |— ccl_draw_colors Generate a color image output

```

6 Questions

Q1 (0pts) : Describe the machine you run this project on: how many physical/logical cores? What system environment? (e.g. *Linux natif, WSL/Windows10, Mac, Machine virtuelle/Windows 8, etc.*)

Q2 (3pts) : Identify the following variables between the algorithm above, and the code. Explain their role.

Symbol in algorithm	Variable name in code	Role (explain in 1 sentence)
I		
E		
T		
N		
n_E^{max}		
n_E		
n_C		

Q3. What functions take most time to run?

Explain, using drawing or diagrams if applicable, how you will parallelize these functions.

Q4: Compile in performance mode, and measure the execution times on image "large.bpm", using 1 to 20 threads, and draw the curve $S(n) = \frac{Real(1)}{Real(n)}$.

Q5: Prepare an oral presentation, 10min long; in English, covering at least:

- What challenges you faced and how you solved them.
- Your results, in terms of performance & speedup,
- Try and illustrate with aesthetic examples 😊

For submitting, please insert your names into the Makefile, then use the `make submit` command to compress your program and report.